

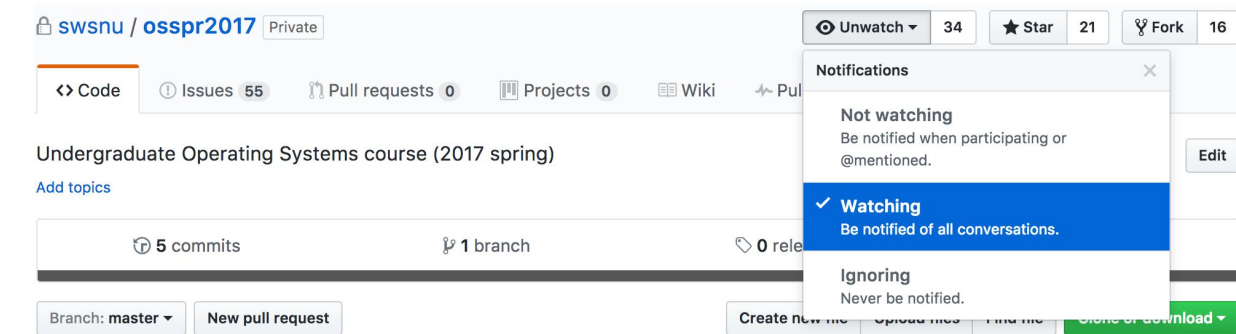
Project 2



CMSLab, SNU

Notification Settings

- <https://github.com/swsnu/osspr2017>
- <https://github.com/settings/emails>
- <https://github.com/settings/notifications>



Email notification preferences

Primary email address

gyeonginyu@gmail.com

Choose which email updates you receive on conversations you're participating in or watching

- ☒ Comments on Issues and Pull Requests
- ☒ Pull Request reviews
- ☒ Pull Request pushes
- ☐ Include your own updates

Project 2 Overview

- Implement rotation-based read/write lock
- Each lock has a “rotation range”
 - Every lock can be acquired when the current rotation is in the rotation range
 - Or, it is blocked until the current rotation is located in the rotation range
- Read lock could be acquired when no acquired **write** lock range is overlapping with its range
- Write lock could be acquired when no acquired **read/write** lock range is overlapping with its range
 - Exclusive access

Rotation Range

- 1 axis: rotation
 - Actually, Tizen orientation has three axes! (Azimuth, Pitch, Roll)
- $(\text{degree} - \text{range}) \leq \mathbf{range} \leq (\text{degree} + \text{range})$
- Rotation ranges are **inclusive**
 - Ex) [30, 60], [60, 90] are overlapped
- Rotation ranges are **circular**
 - Ex) [330, 30] and [30, 330]
 - [330 ... 0 ... 30]
 - [30 ... 180 ... 330]

Range Example (1)

- Rotation 1
 - degree = 30
 - range = 30
- Rotation 2
 - degree = 45
 - range = 30

Are they overlapped?

Range Example (1)

- Rotation 1
 - degree = 30
 - range = 30
- Rotation 1
 - [0, 60]
- Rotation 2
 - degree = 45
 - range = 30
- Rotation 2
 - [15, 75]

Are they overlapped?

Yes!

Range Example (2)

- Rotation 1
 - degree = 30
 - range = 60
- Rotation 2
 - degree = 315
 - range = 30

Are they overlapped?

Range Example (2)


- Rotation 1
 - degree = 30
 - range = 60
- Rotation 1
 - $[0, 90] + [330, 360)$
- Rotation 2
 - degree = 315
 - range = 30
- Rotation 2
 - $[285, 345]$

Are they overlapped?

Yes!

Rotation Lock Example


rotlock_write(3-10)



Rotation Lock Example

rotlock_write(1-4)

device
rotation




success

wait



Rotation Lock Example

rotlock_read(6-9)




 writer
 reader

Rotation Lock Example

rotlock_read(6-8)


device
rotation



 writer
 reader

Rotation Lock Example

rotunlock_write(3-10)




 writer
 reader

Rotation Lock Example

rotunlock_write(3-10)

device
rotation




wait
success
success

 writer
 reader

Rotation Lock Example

Device rotation changed to 3



 writer
 reader


Preventing writer starvation

- You should implement a policy for preventing starvation of writers
 - Why? If reader comes in infinitely, a writer could wait forever!
- **If a reader holds a lock and a writer wants to take the lock, no more readers can take the lock**
- If you design your own *additional* policy, explain that in your README.md file, report, and slides!

How to prevent writer starvation?


- A write lock is waiting for the rotation changes & the reader to release its lock

Acquired locks



Current Rotation = 40


Waiting locks



How to prevent writer starvation?


- A read lock [40, 60] came

Acquired locks



Current Rotation = 40


Waiting locks



How to prevent writer starvation?


- $40 \in [40, 60] \rightarrow$ Acquires its lock immediately

Acquired locks



Current Rotation = 40

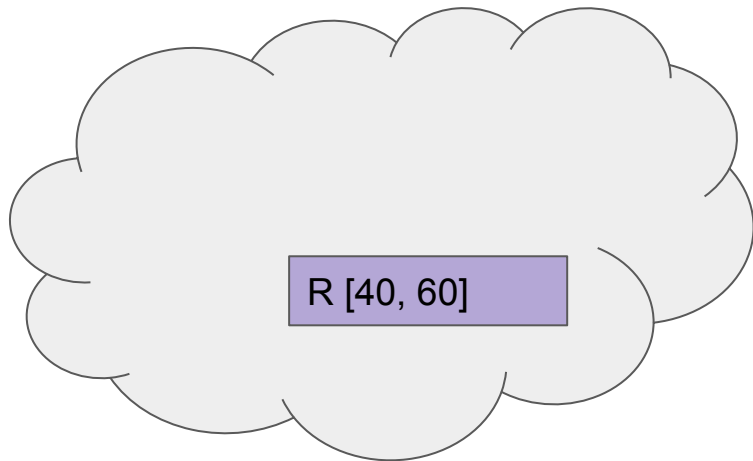
Waiting locks



How to prevent writer starvation?

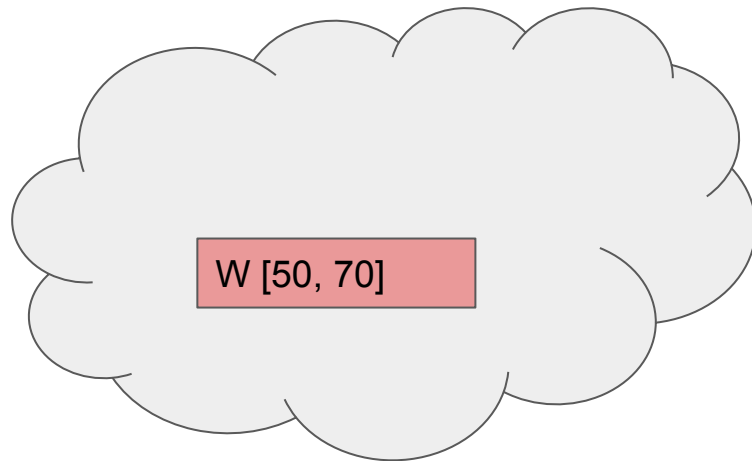
- R [0, 60] releases its lock

Acquired locks



Current Rotation = 40

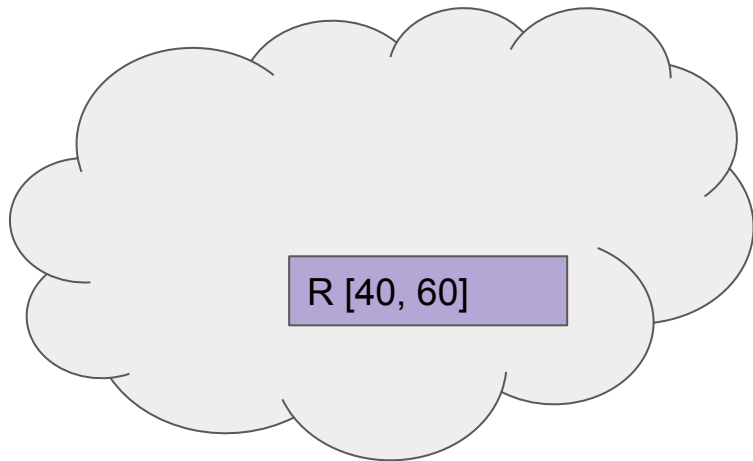
Waiting locks



How to prevent writer starvation?


- Rotation changes $40 \rightarrow 50$

Acquired locks



Current Rotation = 50

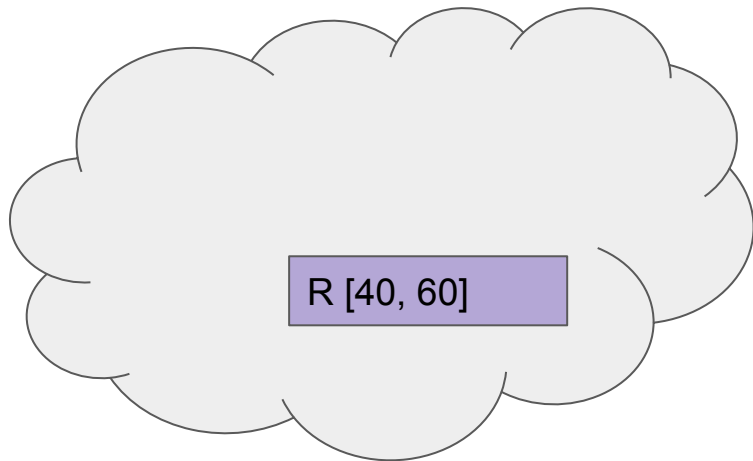
Waiting locks



How to prevent writer starvation?


- A read lock [45, 55] came
- $50 \in [45, 55]$ but cannot acquire its lock, because ...

Acquired locks



Current Rotation = 50


Waiting locks



How to prevent writer starvation?


- W [50, 70] is waiting and $50 \in [50, 70]$
- W [50, 70] is waiting for R [40, 60] to release its lock
- [45, 55] overlaps with [50, 70]

Acquired locks



Current Rotation = 50

Waiting locks



How to prevent writer starvation?

- W [50, 70]
 - W [50, 70]
 - [45, 55] o
- If a reader holds a lock
and a writer wants to take
the lock,
no more readers
can take the lock

R [40, 60]


R [45, 55]

W [50, 70]

Current Rotation = 50

How to prevent writer starvation?

- W [50, 70]
 - W [50, 70]
 - [45, 55] o
- If a reader(R[40, 60]) holds a lock
and a writer(W[50, 70]) wants to take($50 \in [50, 70]$)
the lock([40, 60] overlaps with [50, 70], waiting),
no more readers(R[45, 55])
can take the lock([45, 55] overlaps with [50, 70])




Current Rotation = 50

How to prevent writer starvation?


- Rotation changes 50 \rightarrow 45

Acquired locks



Current Rotation = 45


Waiting locks




How to prevent writer starvation?

- W [50, 70] cannot grab its lock anymore ($\because 45 \notin [50, 70]$)
- Starvation prevention policy is no more applied \rightarrow R [45, 55] acquires its lock

Acquired locks



Waiting locks




Current Rotation = 45

Terminating routine

- When a thread that has some holding or waiting locks is terminating, the remaining locks should be released (holding) or removed (waiting).
- Hints
 - `exit_rotlock()` in `kernel/rotation.c`
 - Release holding locks
 - Remove waiting locks
 - Inject `exit_rotlock()` to `do_exit()` in `kernel/exit.c`

Isolation


- Multiple processes shares same rotation lock system.
 - You have to identify which process (thread) the lock belongs to.
- A process can't release locks that other processes hold.



Isolation

- A process can't release locks that other processes hold.


p0 calls unlock(6-8)



Isolation

- A process can't release locks that other processes hold.

p0 calls unlock(6-8) **again**



Some hints for doing your project

- You will create things like acquired lock list and waiting lock list
 - Different approaches are also possible :)
 - Ex) Bitmap, linked list, ...
- Accessing those lists could result in race conditions
 - Remind: ARTIK has 4 cores
 - You should carefully design your code to prevent possible race conditions
 - Ex) One thread is removing a lock from waiting list, but another thread can access to the waiting list at the same time

Some hints for doing your project

- Possible approaches
 - Global
 - Ex) Manage acquired lock list and waiting lock list using a same lock
 - Fine-grained
 - Ex) Manage acquired lock list and waiting lock list using separated locks
 - Better concurrency, more complicated :)
- Synchronization mechanism
 - Spin Lock
 - Eligible for short sleep (e.g. short list iteration)
 - RCU (Read-Copy-Update)
 - ...

Some hints for doing your project

- How to block processes
 - `set_current_state(TASK_INTERRUPTIBLE)`
 - Wait queue
 - `prepare_to_wait()` and `finish_wait()`
 - wait event
 - Maybe it's necessary to modify existing kernel code :)
or poor performance...
 - condition variable, mutex
 - ...
- **You should call “`schedule()`” right after `set_current_state()`, `prepare_to_wait()`, ...**

Some hints for doing your project

- How to wake up blocked processes?
 - PID or task_struct
 - `wake_up_process(task_struct *t)`
 - You can get “task_struct” value of the blocked process using pid.
`pid_task()`, `find_vpid()`, ...
 - wait queue
 - `wake_up_interruptible()`
 - ...
 - ...

Some hints for doing your project

- Lists could be changed during iteration
 - `list_for_each_entry_safe` could be useful to you
- Please remember that ...
 - The rotation range is **circular!**
 - You should implement a logic for determining two circular ranges are overlapping or not
 - Be aware of deadlocks!

Wait queue

- Blocking
 - Declaring a new wait queue
 - `DECLARE_WAIT_QUEUE_HEAD()`
 - Create wait queue entry.
 - `DEFINE_WAIT()`, `DECLARE_WAITQUEUE()`, `DEFINE_WAIT_FUNC()`
 - Register a function that wakes up with the entry.
 - The function can contain condition check code.
 - `DEFINE_WAIT()` registers `autoremove_wake_function` automatically.
 - Then, `prepare_to_wait()` and `schedule()`.

Wait queue

- Wake up
 - Other processes call `__wake_up`.
 - `__wake_up` calls the function registered with the entry.
 - If the function doesn't remove the entry from the wait queue, the woken process handles it.
 - `finish_wait()`

Wait queue

- wait_event
 - DEFINE_WAIT()
 - while(true) {
 - prepare_to_wait()
 - if (condition)
 - break;
 - schedule();
 - }
 - finish_wait()

Condition variable vs Wait queue

- Condition variable can be implemented with wait queue.
 - wait()
 - Register a function that check condition with the wait queue entry.
 - DEFINE_WAIT_FUNC(name, your_function)
 - Or use wait_event.
 - signal(): __wake_up(..., nr_exclusive=1, ...)
 - broadcast() : __wake_up(..., nr_exclusive=0, ...)

Selector & Trial

- Selector & Trial require a same lock (**$0 \leq \text{range} \leq 180$**)
 - If current rotation is 240, both selector & trial wait.
- When the device rotation is out of that range, both Selector & Trial stop working
- When the device rotation gets inside that range, Selector & Trial start to work

Selector & Trial

Selector	Trial
write_lock	
.	read_lock & wait
10	
write_unlock	acquire lock
write_lock & wait	
	10 = 2 * 5
acquire lock	read_unlock

Q & A